

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:

```
traindf=pd.read_csv(r"C:\Users\shaha\OneDrive\Desktop\Excel\Rainfall train.csv")
traindf
```

Out[3]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.0
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.0
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.0
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.0
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.0
...
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527.0
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636.0
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352.0
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592.0
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217.0

41 rows × 19 columns



In [4]:

```
testdf=pd.read_csv(r"C:\Users\shaha\OneDrive\Desktop\Excel\rainfall in india.csv")
testdf
```

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	O
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	38
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	19
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	18
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	22
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	26
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	11
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	14
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	7
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	16
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	16

4116 rows × 19 columns



In [5]:

```
traindf.head()
```

Out[5]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP
0	ANDAMAN And NICOBAR ISLANDS	NICOBAR	107.3	57.9	65.2	117.0	358.5	295.5	285.0	271.9	354.8
1	ANDAMAN And NICOBAR ISLANDS	SOUTH ANDAMAN	43.7	26.0	18.6	90.5	374.4	457.2	421.3	423.1	455.6
2	ANDAMAN And NICOBAR ISLANDS	N & M ANDAMAN	32.7	15.9	8.6	53.4	343.6	503.3	465.4	460.9	454.8
3	ARUNACHAL PRADESH	LOHIT	42.2	80.8	176.4	358.5	306.4	447.0	660.1	427.8	313.6
4	ARUNACHAL PRADESH	EAST SIANG	33.3	79.5	105.9	216.5	323.0	738.3	990.9	711.2	568.0

In [6]:

```
testdf.head()
```

Out[6]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	3
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	3
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	

In [7]:

```
traindf.tail()
```

Out[7]:

	STATE_UT_NAME	DISTRICT	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG
636	KERALA	IDUKKI	13.4	22.1	43.6	150.4	232.6	651.6	788.9	527
637	KERALA	KASARGOD	2.3	1.0	8.4	46.9	217.6	999.6	1108.5	636
638	KERALA	PATHANAMTHITTA	19.8	45.2	73.9	184.9	294.7	556.9	539.9	352
639	KERALA	WAYANAD	4.8	8.3	17.5	83.3	174.6	698.1	1110.4	592
640	LAKSHADWEEP	LAKSHADWEEP	20.8	14.7	11.8	48.9	171.7	330.2	287.7	217

In [8]:

```
testdf.tail()
```

Out[8]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.3
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4

In [9]:

```
traindf.shape
```

Out[9]:

(641, 19)

In [10]:

```
testdf.shape
```

Out[10]:

(4116, 19)

In [11]:

```
traindf.describe()
```

Out[11]:

	JAN	FEB	MAR	APR	MAY	JUN	JUL
count	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000	641.000000
mean	18.355070	20.984399	30.034789	45.543214	81.535101	196.007332	326.033697
std	21.082806	27.729596	45.451082	71.556279	111.960390	196.556284	221.364643
min	0.000000	0.000000	0.000000	0.000000	0.900000	3.800000	11.600000
25%	6.900000	7.000000	7.000000	5.000000	12.100000	68.800000	206.400000
50%	13.300000	12.300000	12.700000	15.100000	33.900000	131.900000	293.700000
75%	19.200000	24.100000	33.200000	48.300000	91.900000	226.600000	374.800000
max	144.500000	229.600000	367.900000	554.400000	733.700000	1476.200000	1820.900000

In [12]:

```
testdf.describe()
```

Out[12]:

	YEAR	JAN	FEB	MAR	APR	MAY	
count	4116.000000	4112.000000	4113.000000	4110.000000	4112.000000	4113.000000	4111.000000
mean	1958.218659	18.957320	21.805325	27.359197	43.127432	85.745417	230.218659
std	33.140898	33.585371	35.909488	46.959424	67.831168	123.234904	234.718659
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.418659
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.318659
50%	1958.000000	6.000000	6.700000	7.800000	15.700000	36.600000	138.718659
75%	1987.000000	22.200000	26.800000	31.300000	49.950000	97.200000	305.118659
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.918659

In [13]:

```
traindf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 641 entries, 0 to 640
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   STATE_UT_NAME         641 non-null    object
1   DISTRICT              641 non-null    object
2   JAN                   641 non-null    float64
3   FEB                   641 non-null    float64
4   MAR                   641 non-null    float64
5   APR                   641 non-null    float64
6   MAY                   641 non-null    float64
7   JUN                   641 non-null    float64
8   JUL                   641 non-null    float64
9   AUG                   641 non-null    float64
10  SEP                   641 non-null    float64
11  OCT                   641 non-null    float64
12  NOV                   641 non-null    float64
13  DEC                   641 non-null    float64
14  ANNUAL                641 non-null    float64
15  Jan-Feb               641 non-null    float64
16  Mar-May               641 non-null    float64
17  Jun-Sep               641 non-null    float64
18  Oct-Dec               641 non-null    float64
dtypes: float64(17), object(2)
memory usage: 95.3+ KB
```

In [14]:

testdf.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION     4116 non-null   object
1   YEAR            4116 non-null   int64
2   JAN             4112 non-null   float64
3   FEB             4113 non-null   float64
4   MAR             4110 non-null   float64
5   APR             4112 non-null   float64
6   MAY             4113 non-null   float64
7   JUN             4111 non-null   float64
8   JUL             4109 non-null   float64
9   AUG             4112 non-null   float64
10  SEP             4110 non-null   float64
11  OCT             4109 non-null   float64
12  NOV             4105 non-null   float64
13  DEC             4106 non-null   float64
14  ANNUAL          4090 non-null   float64
15  Jan-Feb        4110 non-null   float64
16  Mar-May        4107 non-null   float64
17  Jun-Sep        4106 non-null   float64
18  Oct-Dec        4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [15]:

traindf.isnull().sum()

Out[15]:

```
STATE_UT_NAME    0
DISTRICT         0
JAN              0
FEB              0
MAR              0
APR              0
MAY              0
JUN              0
JUL              0
AUG              0
SEP              0
OCT              0
NOV              0
DEC              0
ANNUAL           0
Jan-Feb          0
Mar-May          0
Jun-Sep          0
Oct-Dec          0
dtype: int64
```

In [16]:

```
testdf.isnull().sum()
```

Out[16]:

```
SUBDIVISION    0
YEAR           0
JAN            4
FEB            3
MAR            6
APR            4
MAY            3
JUN            5
JUL            7
AUG            4
SEP            6
OCT            7
NOV           11
DEC           10
ANNUAL         26
Jan-Feb        6
Mar-May        9
Jun-Sep       10
Oct-Dec       13
dtype: int64
```

In [17]:

```
testdf.dropna(inplace=True)
```

In [18]:

```
testdf.isnull().sum()
```

Out[18]:

```
SUBDIVISION    0
YEAR           0
JAN            0
FEB            0
MAR            0
APR            0
MAY            0
JUN            0
JUL            0
AUG            0
SEP            0
OCT            0
NOV            0
DEC            0
ANNUAL         0
Jan-Feb        0
Mar-May        0
Jun-Sep        0
Oct-Dec        0
dtype: int64
```


In [26]:

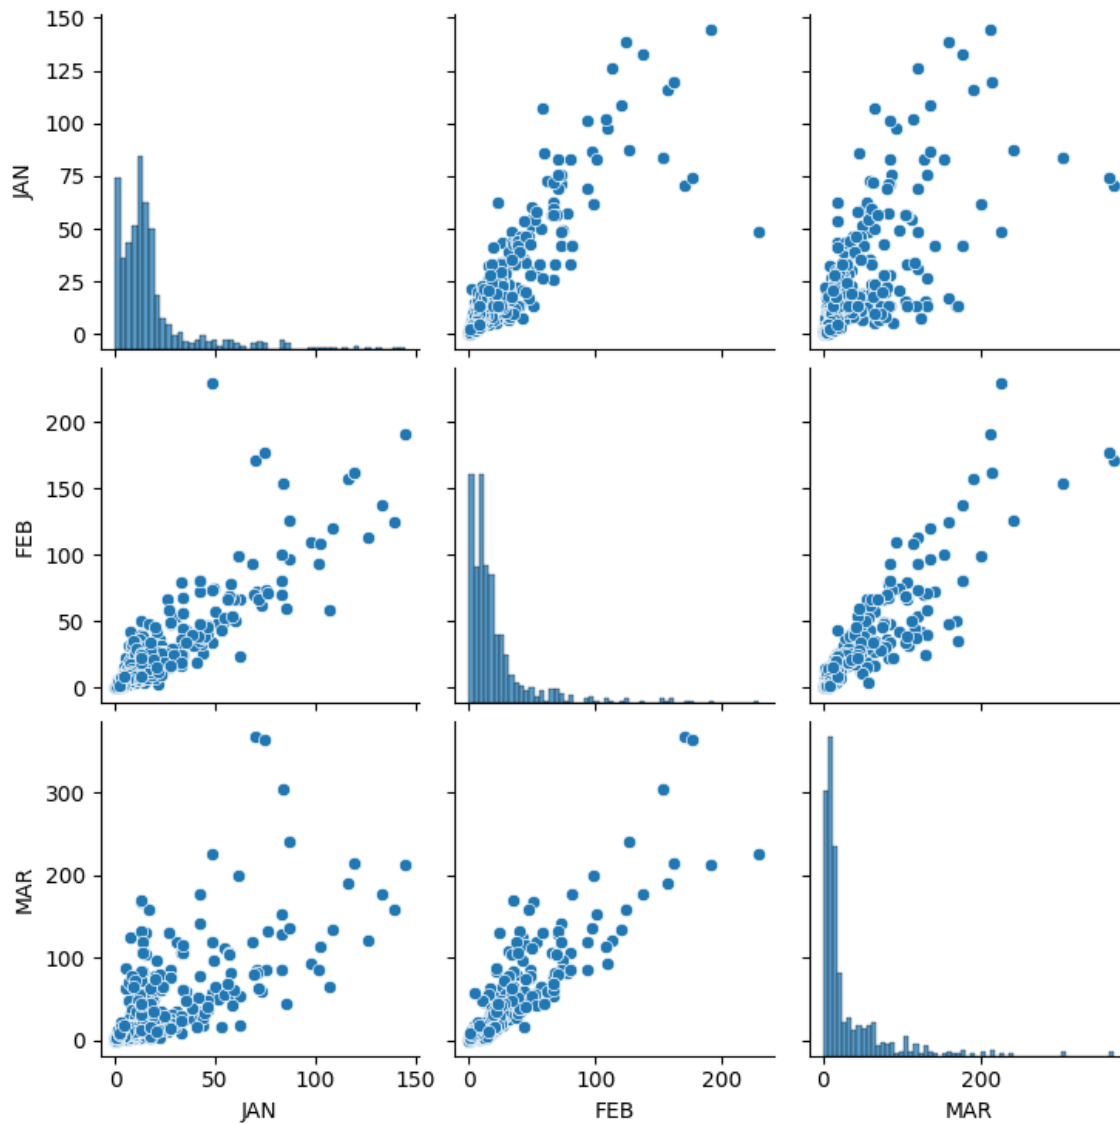
```
data=traindf[['JAN', 'FEB', 'MAR']]
```

In [28]:

```
sns.pairplot(data)
```

Out[28]:

<seaborn.axisgrid.PairGrid at 0x2cd56029db0>



In [29]:

```
traindf['STATE_UT_NAME'].value_counts()
```

Out[29]:

STATE_UT_NAME	
UTTAR PRADESH	71
MADHYA PRADESH	50
BIHAR	38
MAHARASHTRA	35
RAJASTHAN	33
TAMIL NADU	32
KARNATAKA	30
ORISSA	30
ASSAM	27
GUJARAT	26
JHARKHAND	24
ANDHRA PRADESH	23
JAMMU AND KASHMIR	22
HARYANA	21
PUNJAB	20
WEST BENGAL	19
CHATISGARH	18
ARUNACHAL PRADESH	16
KERALA	14
UTTARANCHAL	13
HIMACHAL	12
NAGALAND	11
MIZORAM	9
MANIPUR	9
DELHI	9
MEGHALAYA	7
SIKKIM	4
TRIPURA	4
PONDICHERRY	4
ANDAMAN And NICOBAR ISLANDS	3
GOA	2
DAMAN AND DIU	2
DADAR NAGAR HAVELI	1
CHANDIGARH	1
LAKSHADWEEP	1

Name: count, dtype: int64

In [30]:

```
traindf['DISTRICT'].value_counts()
```

Out[30]:

```
DISTRICT
BIJAPUR      2
BILASPUR     2
AURANGABAD   2
HAMIRPUR     2
NICOBAR      1
..
GONDA        1
GORAKHPUR    1
HARDOI       1
JAUNPUR      1
LAKSHADWEEP  1
Name: count, Length: 637, dtype: int64
```

Linear Regression

In [31]:

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

In [40]:


```
x=traindf[['JAN','FEB','MAR','APR','MAY','JUN']]
y=traindf['ANNUAL']
```

In [41]:

```
sns.heatmap(x.corr())
```

Out[41]:

<Axes: >



In [42]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=100)
```

In [43]:

```
regr = LinearRegression()
regr.fit(x_train,y_train)
print(regr.intercept_)
coeff_train_df=pd.DataFrame(regr.coef_,x.columns,columns=['coefficient'])
coeff_train_df
```

437.52725826332585

Out[43]:

	coefficient
JAN	3.017735
FEB	6.071487
MAR	-3.029124
APR	0.633801
MAY	0.825779
JUN	3.683982

In [44]:

```
score=regr.score(x_test,y_test)
print(score)
```


0.9488529666589274

In [45]:

```
predictions=regr.predict(x_test)
plt.scatter(y_test,predictions)
```

Out[45]:

<matplotlib.collections.PathCollection at 0x2cd6e870430>



In [49]:

```
x=np.array(traindf['JAN']).reshape(-1,1)
y=np.array(traindf['ANNUAL']).reshape(-1,1)
```

In [50]:


```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(x_train,y_train)
regr.fit(x_test,y_test)
```

Out[50]:

```
LinearRegression
LinearRegression()
```

In [51]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='b')
plt.show()
```



Ridge and Lasso regression

In [59]:

```
#Ridge Regression MOdel
from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

In [60]:

```
plt.figure(figsize=(10,10))
```

Out[60]:

<Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

In [115]:

```
features = traindf.columns[2:5]
target = traindf.columns[-1:]
#x and y values
x = traindf[features].values
y = traindf[target].values
#split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
#Scale features
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of x_train is (448, 3)
The dimension of x_test is (193, 3)

In [116]:

```
#Model
lr = LinearRegression()
#fit model
lr.fit(x_train ,y_train)
#predict
#prediction = lr.predict(x_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.08412859673411976
The test score for lr model is 0.03895606778221439

In [117]:

```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train,y_train)
test_score_ridge = ridgeReg.score(x_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.08354804861559051
The test score for ridge model is 0.03961593206840308

In [118]:

```
#Lasso Regression model
print("\nLasso Model:\n")
lasso = Lasso(alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls = lasso.score(x_train,y_train)
test_score_ls = lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:


The train score for ls model is 0.07504785033558381
The test score for ls model is 0.052588033859716776

In [119]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "pie")
```

Out[119]:

<Axes: >



In [120]:

```
#using the linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
#score
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

0.07504785033558381
0.052588033859716776

In [122]:

```
#using the linear Cv model
from sklearn.linear_model import RidgeCV
#Ridge cross validation
ridge_cv = RidgeCV(alphas = [0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(x_train,y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(x_test,y_test)))
```

The train score for ridge model is 0.08354804861559062
The train score for ridge model is 0.039615932068407744

Elastic Net

In [123]:

```
from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
regr.score(x,y)
```

[0.25536298 -1.31993694 1.50300354]
[120.58314834]

Out[123]:

0.07698323574326038

In [124]:

```
y_pred_elastic = regr.predict(x_train)
```

In [125]:

```
mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
print("Mean squared Error on test set",mean_squared_error)
```

Mean squared Error on test set 22500.582660635675

conclusion

from the given Rainfall dataset,we have performed LinearRegression model,Ridge and Lasso Regression,Elastic Net.After applying the models,Linear Regression gives the best score:0.9488529666589274.so,Linear Regression model is the Best fit.

In []: